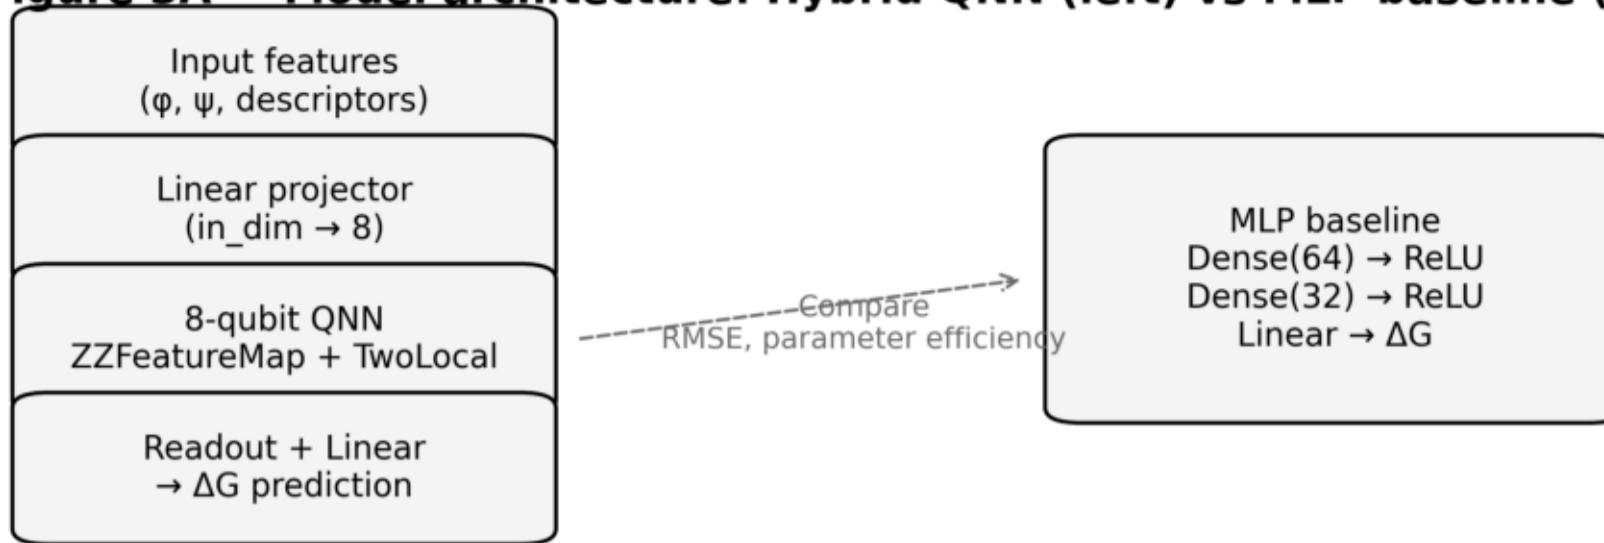


Figure 3A — Model Architecture

Figure 3A — Model architecture: Hybrid QNN (left) vs MLP baseline (right)



Hybrid QNN: projector \rightarrow n-qubit variational circuit (ZZFeatureMap + TwoLocal) \rightarrow readout

Figure 3B — Parity Plot

Figure 3B — Parity: Pred vs Experimental tre

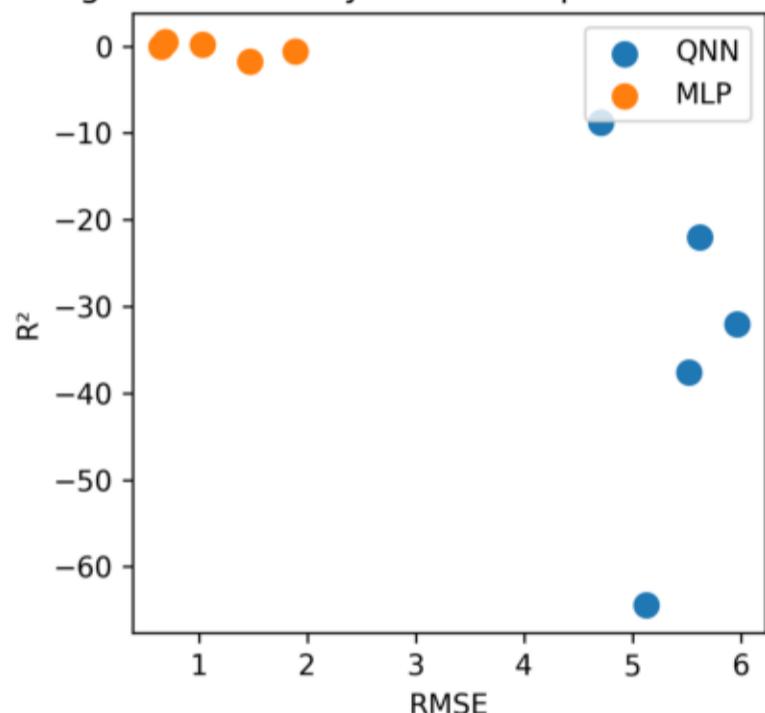


Figure 3C — Training Convergence

Figure 3C — Training Convergence

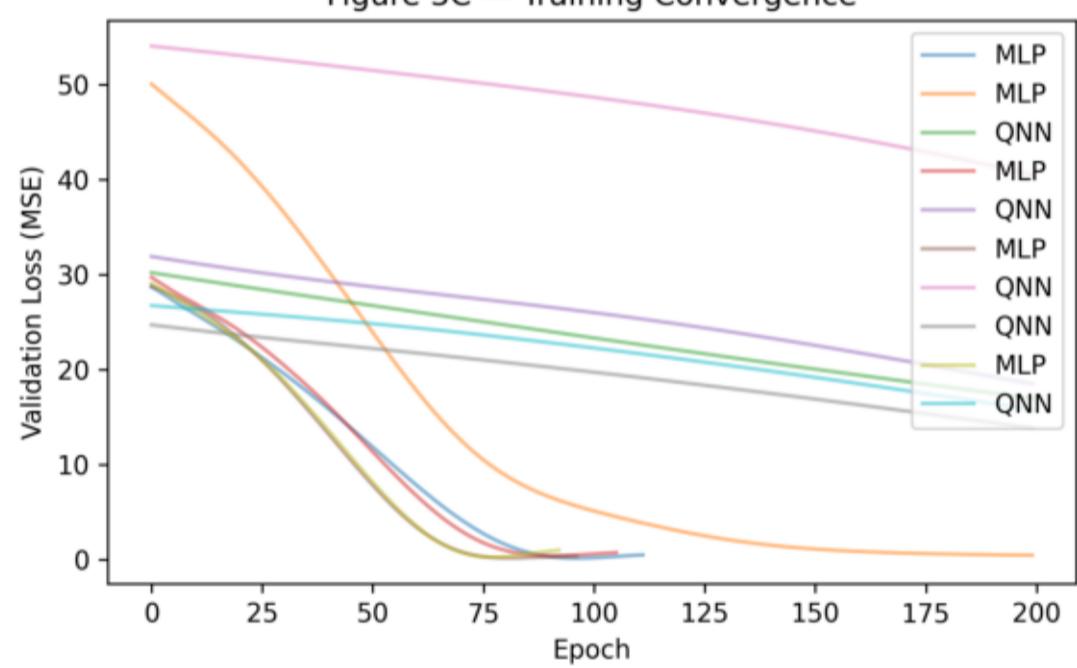


Figure 3D — Parameter Efficiency

Figure 3D — Parameter Efficiency

